← Back Final Quiz
Graded Quiz * 1 min

	Congratulations! You passed! Grade Latest Submission To pass 80% or received 95% Grade 90% higher	o to next item
1.	Which factor drives the data model in a NoSQL database?	1/1 point
	 The schema The queries you intend to make Relationships among entities 	
	The complexity of the joins ∠ Expand	
	Correct Correct! In NoSQL it's not the data that drives your data model or schema. It's the way your application accesses the data and the queries you are going to make.	
2.	Which of the following is a characteristic of a NoSQL database? Tables are related using foreign keys.	1/1 point
	 System capacity is easily scaled. Data is stored on a single node. Similar data is stored in the same location. 	
	∠ Expand ✓ Correct	
3.	Correct! You can easily or scale to increase system capacity, just by adding new servers to the cluster. In the ACID model what does it mean for the data to be "durable"?	0 / 1 point
	 Transactions cannot compromise the integrity of other transactions by interacting with them while they are still in progress. All operations in a transaction succeed or every operation is rolled back. On the completion of a transaction, the structural integrity of the data in the database is not compromised. 	
	☐ If a transaction fails, it will not impact the already changed data. ☑ Expand	
	Ncorrect Incorrect. Review the ACID vs BASE video.	
4.	Which of the following makes NoSQL databases most appropriate for use with cloud computing? They support complex queries. They are denormalized.	1/1 point
	 Their distributed nature makes them easily deployed and operated on clusters of servers. They have intuitive data structures. 	
	Correct Correct! The distributed data nature of NoSQL databases means that they can be deployed and operated on clusters of servers in cloud architectures, thereby massively reducing cost."	
5.	Which of the following statements best describes NoSQL? NoSQL is a family of open source relational databases	1/1 point
	NoSQL is a family of open source non-relational databases that differ greatly in style and technology NoSQL is a family of open source relational databases that do not include SQL NoSQL is a family of open source non-relational databases that share exactly the same technologies	
	Expand	
	Correct The term NoSQL refers to a family of databases that vary widely in style and technology, but which all share a common trait in that they are non-relational in nature. Contrary to what it sounds like, 'NoSQL' actually stands for 'Not Only SQL', not 'NO SQL'.	
6.	In MongoDB, which of the following common aggregation stages takes the outcome from the previous stage and stores it in a target collection? • Smerge	1/1 point
	\$count\$sort\$project	
	 ✓ Expand ✓ Correct The "\$merge" aggregation stage takes the outcome from the previous stage and stores it in a target 	
7.	Collection. True or False: Since the address format can change from one country to another Mongo would be more	1/1 point
	 appropriate to store addresses from multiple countries than a relational database. True False 	
	∠ [¬] Expand	
	Correct Correct! Mongo allows us to store unstructured data.	
8.	Which statement about MongoDB is true? Data in the same document must be of the same data type. Documents are linked by primary keys.	1/1 point
	 It requires third-party software for complex queries. You can design complex data structures easily. 	
	 ✓ Correct Correct! You can design complex data structures easily in MongoDB without worrying about the complexity of how it is stored and how it should be linked. 	
9.	Which everyday example is similar to indexing in MongoDB?	1/1 point
	A blueprint A table of contents A recipe A taxonomy	
	Expand	
10	Correct! Indexes are key-value pairs sorted in ascending order. D. What is the primary purpose of replication?	1/1 point
	 If one server fails, you still have multiple copies of the data present. It prevents you from accidentally deleting data. It increases throughput by directing queries to relevant partitions. 	
	○ It allows you to partition large data sets. ∠ Expand	
	Correct Correct! It provides you with a highly available database during failures or during periods of planned maintenance.	
11	What would you most likely use blobs for in Cassandra Query Language (CQL)? Storing a high range of integers Representing encoded strings	1/1 point
	Storing key-value pairs Storing multimedia objects	
	 ∠ Expand ✓ Correct Blobs are typically used to store images, audio, or other multimedia objects. 	
12	. How does the syntax of Cassandra Query Language (CQL) support lightweight transactions? Using a WHERE clause within READ and DELETE statements.	1/1 point
	Using an IF clause within READ and DELETE statements. Using a WHERE clause within INSERT and UPDATE statements. Using an IF clause within INSERT and UPDATE statements.	
	∠ Expand ✓ Correct	
13	Correct! Lightweight Transactions are supported by introducing IF in INSERT and UPDATE statements." 3. Which features set Cassandra Apache apart from other NoSQL databases?	1/1 point
	 Primary/secondary architecture and fast write throughput Simple peer-to-peer architecture and scalability Primary/secondary architecture and data consistency Simple peer-to-peer architecture and fast write throughput 	
	∠ [¬] Expand	
	Correct! Cassandra has a simple peer-to-peer architecture. It also maintains cluster performance for read capability to scale clusters fast, making it useful for services that need to be accessed worldwide.	
14	What is the purpose of the partition key hash? To determine the location of the data in a cluster. To optimize the read performance of queries.	1/1 point
	 It is used to determine what node and subsequent replicas will get the data. To add uniqueness to each entry in a table. 	
	Correct Correct! The partition key hash is used to determine what node (and subsequent replicas) will get the data.	
15	5. What do you need to start data distribution in Cassandra? O Data sets	0 / 1 point
	Opata sets Queries A key value hash Tokens	
	Expand Note: The second secon	
16	Review the Key Features of Apache Cassandra video. 6. When you create a new IBM Cloudant database, what do you need to select?	1/1 point
	 The number of documents The replication type The curl command line tool The partitioning type 	
	∠ [¬] Expand	
	Correct When you create a database, you must select whether you want the database to be partitioned or non-partitioned.	
17	What is the default HTTP method used in curl if the method is not specified? PUT GET	1/1 point
	O DELETE O POST	
	Correct Correct! Curl defaults to using the 'GET' HTTP method unless you specify another method using the "-X" command-line switch.	
18	3. How are IBM Cloudant documents stored?	1/1 point
	As files As tables As objects As JSON	
	✓ Expand ✓ Correct	
19	Correct! With Cloudant, documents are stored in the popular JSON format with a flexible schema. O. Which statement best describes IBM Cloudant?	1 / 1 point
	A data warehouse A document-based NoSQL database A relational database	
	○ A hierarchical database ∠ Expand	
	✓ Correct Correct! Cloudant is a document-based NoSQL database, it is not a data warehouse, like Hadoop is.	
20	What software is required to work with IBM Cloudant's databases? Proprietary software is not required. Python 2.6 or higher	1/1 point
	IBM CloudBuilder IBM Cloudant Integrated Development Environment (IDE)	
	 Expand Correct Correct! A device does not require specialized software, or drivers, or proprietary protocols, or anything of that nature, in order to work with Cloudant's databases. 	